Chen Jiawei

AI Infra Researcher & Engineer

2451427796@qq.com Beijing, China cjwystz

Profile

AI Infra researcher and engineer focused on large-model distributed training & inference, memory optimization and load scheduling. Open-source contributor to FlagOS and Hugging Face PEFT with 1000-GPU-scale experience on heterogeneous accelerators. Seeking a PhD in large-model systems (2028-2029 Fall).

Education

B.Eng. University of Science & Technology Beijing, Information Security — Beijing, China — 2023 – 2027

- IELTS 6.5 (all >= 6.0)
- Systems security, parallel computing, AI systems
- National First Prize, National Metallurgy Science & Technology Competition
- National Second Prize, ASC Student Supercomputer Challenge

Experience

Qingcheng Jizhi, AI Infra Intern — China — 2025 – present — 1 year

Contribute to the FlagOS / FlagScale training framework for the MetaX ecosystem — bug fixes, hardware adaptation and training-framework optimization.

- 1000-GPU-scale training & inference on Huawei Ascend 910C, MetaX C550, Ali Zhenwu 810E
- FlagScale / Megatron-LM-FL / TransformerEngine-FL: MetaX adaptation and training optimization

FlagOS Open Source Community, Contributor — Remote — 2025 – present — 1 year

Contributions to the FlagOS training stack (te-fl, FlagScale, Megatron-LM adaptations).

- TE package adaptation #128 / FlagScale MetaX fixes #1296
- JIT memory fix #185 / Chunk CE #126
- Qwen3.5 multimodal packing #1308 / MetaX TP overlap adaptation

Hugging Face PEFT, Open-source Contributor — Remote — Oct 2026 – present — 1 month

LoRA-FA optimizer closure gradient-context fix and regression test

- PR #3966 / Technical write-up

Publications

Heterogeneous-modeling Packing Scheduling for General LLM Training — Oct 2026

Packing time-cost model with snake-dealing scheduling; Chunk-CE / Chunk-Linear memory operators; +20% end-to-end training throughput.
Jiawei Chen
(MLSys 2027 (to be submitted, Oct 2026))

Formal-Proof-Guided Multimodal Geometric Reasoning — Jan 2026

GeoGoal benchmark with Skeleton Rate; SGVR dense process reward; +9.7% geometry / +8.0% general math under GRPO.
Jiawei Chen
arxiv.org/abs/2601.05073 (TMLR 2027 (under review))

RAG-Optimized Text-to-Image Generation for Consumer Platforms Sept 2025
RAG-fronted prompt optimization for T2I systems; Tavily + BGE-m3 + ChromaDB two-stage pipeline; +114% average score vs LLM-only rewriting.
Jiawei Chen
github.com/Dadada66666/T2I (WISE 2025 Workshop (published))

Awards

National First Prize: National Metallurgy Science & Technology Competition (2025)

National Second Prize: ASC Student Supercomputer Challenge (2025)

Honorable Mention: MCM / ICM (2024)

Provincial Third Prize: China Collegiate Cyber-Security Challenge (2023)

Skills

Training Systems: Megatron-LM, MindSpeed, FlagScale, FSDP, PyTorch

Inference Systems: vLLM, vLLM-Ascend, SGLang

Hardware Platforms: Huawei Ascend 910C, MetaX C550, Ali Zhenwu 810E, NVIDIA A800/A100

Tooling: Performance Profiler, Docker, AI-assisted coding

Languages

Chinese: Native

English: IELTS 6.5 (all >= 6.0)

Interests

LLM Systems: Distributed training & inference, memory optimization, load scheduling

Embodied AI: Robotics, multimodal reasoning